

Introduction to Database System

ER to Relational



Hi, welcome back.

Why Study the Relational Model?

- Most widely used model.
 - Vendors: IBM, Informix, Microsoft, Oracle, Sybase, etc.

My SQL

PostgreSQL

They all follow the relational model.



Relational Database: Definitions

- *Relational database*: a set of *relations* ← *Table*
- *Relation*: made up of 2 parts:
 - *Instance*: a *table*, with rows and columns.
#Rows = *cardinality*, #fields = *degree / arity*.
 - *Schema*: specifies name of relation, plus name and type of each column.
 - E.G. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).
- Can think of a relation as a *set* of rows or *tuples* (i.e., all rows are distinct).



So, again the relation can be thought of as a set of rows or tuples.

Example Instance of Students Relation

3	sid	name	login	age	gpa
	53666	Jones	jones@cs	18	3.4
	53688	Smith	smith@eecs	18	3.2
	53650	Smith	smith@math	19	3.8

- ❖ Cardinality = 3, degree = 5, all rows distinct
- ❖ Do all columns in a relation instance have to be distinct?



Creating Relations in SQL

- Creates the Students relation. Observe that the type **(domain)** of each field is specified, and enforced by the DBMS whenever tuples are added or modified.
- As another example, the Enrolled table holds information about courses that students take.

CREATE TABLE Students

(sid: CHAR(20),
name: CHAR(20),
login: CHAR(10),
age: INTEGER,
gpa: REAL)

← field
names

CREATE TABLE Enrolled

(sid: CHAR(20),
cid: CHAR(20),
grade: CHAR(2))

So, you can create as many tables using the create table statement.

Destroying and Altering Relations

DROP TABLE Students

- Destroys the relation Students. The schema information *and* the tuples are deleted.

ALTER TABLE Students

ADD COLUMN firstYear: integer

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

even drop it or change it, alter it.



Adding and Deleting Tuples

- Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

- ❖ Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

all the rows for all the people named Smith will be deleted using this command in SQL.



If you would like to insert a new company tuple in your corporation's table, what is the most likely correct command?

☐ INSERT INTO COMPANY (NAME, CITY, STATE)

("ASU CORP", "TEMPE", "AZ")

☒ INSERT INTO COMPANY (NAME, CITY, STATE)

VALUES ("ASU CORP", "TEMPE", "AZ")

Correct

The column names of the table Company and the values for those columns are correctly specified.

☐ INSERT INTO COMPANY

VALUES(3, "ASU CORP", "TEMPE", "AZ")

Continue

Integrity Constraints (ICs)

- **IC**: condition that must be true for *any* instance of the database
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- A legal instance of a relation is one that satisfies all specified ICs.
 - DBMS should not allow illegal instances.
- If the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - Avoids data entry errors, too!



Primary Key Constraints

- A set of fields is a key for a relation if :
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
 - Part 2 false? A *superkey*.
 - If there's >1 key for a relation, one of the keys is chosen (by DBA) to be the *primary key*.
- E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.



multiple keys in the database and you choose one of them to be the primary key.

Primary and Candidate Keys in SQL

- Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the *primary key*.
- ❖ “For a given student and course, there is a single grade.”
vs. “Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade.”
- ❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid,cid) )
```

```
CREATE TABLE Enrolled  
(sid CHAR(20)  
  cid CHAR(20),  
  grade CHAR(2),  
  PRIMARY KEY (sid),  
  UNIQUE (cid, grade) )
```

make that unique as well.



Foreign Keys, Referential Integrity

- Foreign key : Set of fields in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a 'logical pointer'.
- E.g. *sid* is a foreign key referring to **Students**:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.

So, in that case, you can have for example,



Foreign Keys in SQL

- Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled  
(sid CHAR(20), cid CHAR(20), grade CHAR(2),  
PRIMARY KEY (sid,cid),  
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled			Students				
sid	cid	grade	sid	name	login	age	gpa
53666	Carnatic101	C	53666	Jones	jones@cs	18	3.4
53666	Reggae203	B	53688	Smith	smith@eecs	18	3.2
53650	Topology112	A	53650	Smith	smith@math	19	3.8
53666	History105	B					

So, in other words,



Why is referential integrity important?

- ☒ It ensures the correctness of data within a DBMS.

Correct

Any foreign key field must agree with the primary key that is referenced by the foreign key.

- ☐ It ensures a tight relationship between foreign keys in a table.
- ☐ It ensures multiple primary keys to exist in a tuple.

Continue

Enforcing Referential Integrity (2/2)

- What should be done if a Student's tuple is deleted?
 - Also delete all Enrolled tuples that refer to it.
 - ~~Disallow deletion of a Student's tuple that is referred to.~~
 - Set sid in Enrolled tuples that refer to it to a *default sid*.
 - (In SQL, also: Set sid in Enrolled tuples that refer to it to a special value *null*, denoting 'unknown' or 'inapplicable'.)

So, there are several ways to.



Referential Integrity in SQL

- SQL support all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
```

```
(ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

Database systems to give you many options in SQL to define that,

